

# Convolutional Neural Networks

---

Fundamental Concepts & Hands-on Examples

# Slide Deck and Code

[http://bit.ly/cnn\\_ds\\_dphi20](http://bit.ly/cnn_ds_dphi20)

# Session Agenda



Introduction



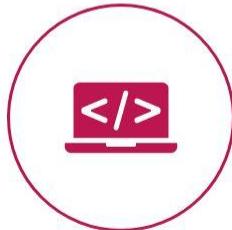
Understanding Convolutional  
Neural Networks



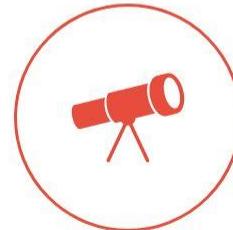
CNN Layer Operations



Transfer Learning Brief



Hands-on Examples



CNN Applications

# Introduction



About Me

# Dipanjan Sarkar

Data Science Lead, Author, Google Developer Expert - ML

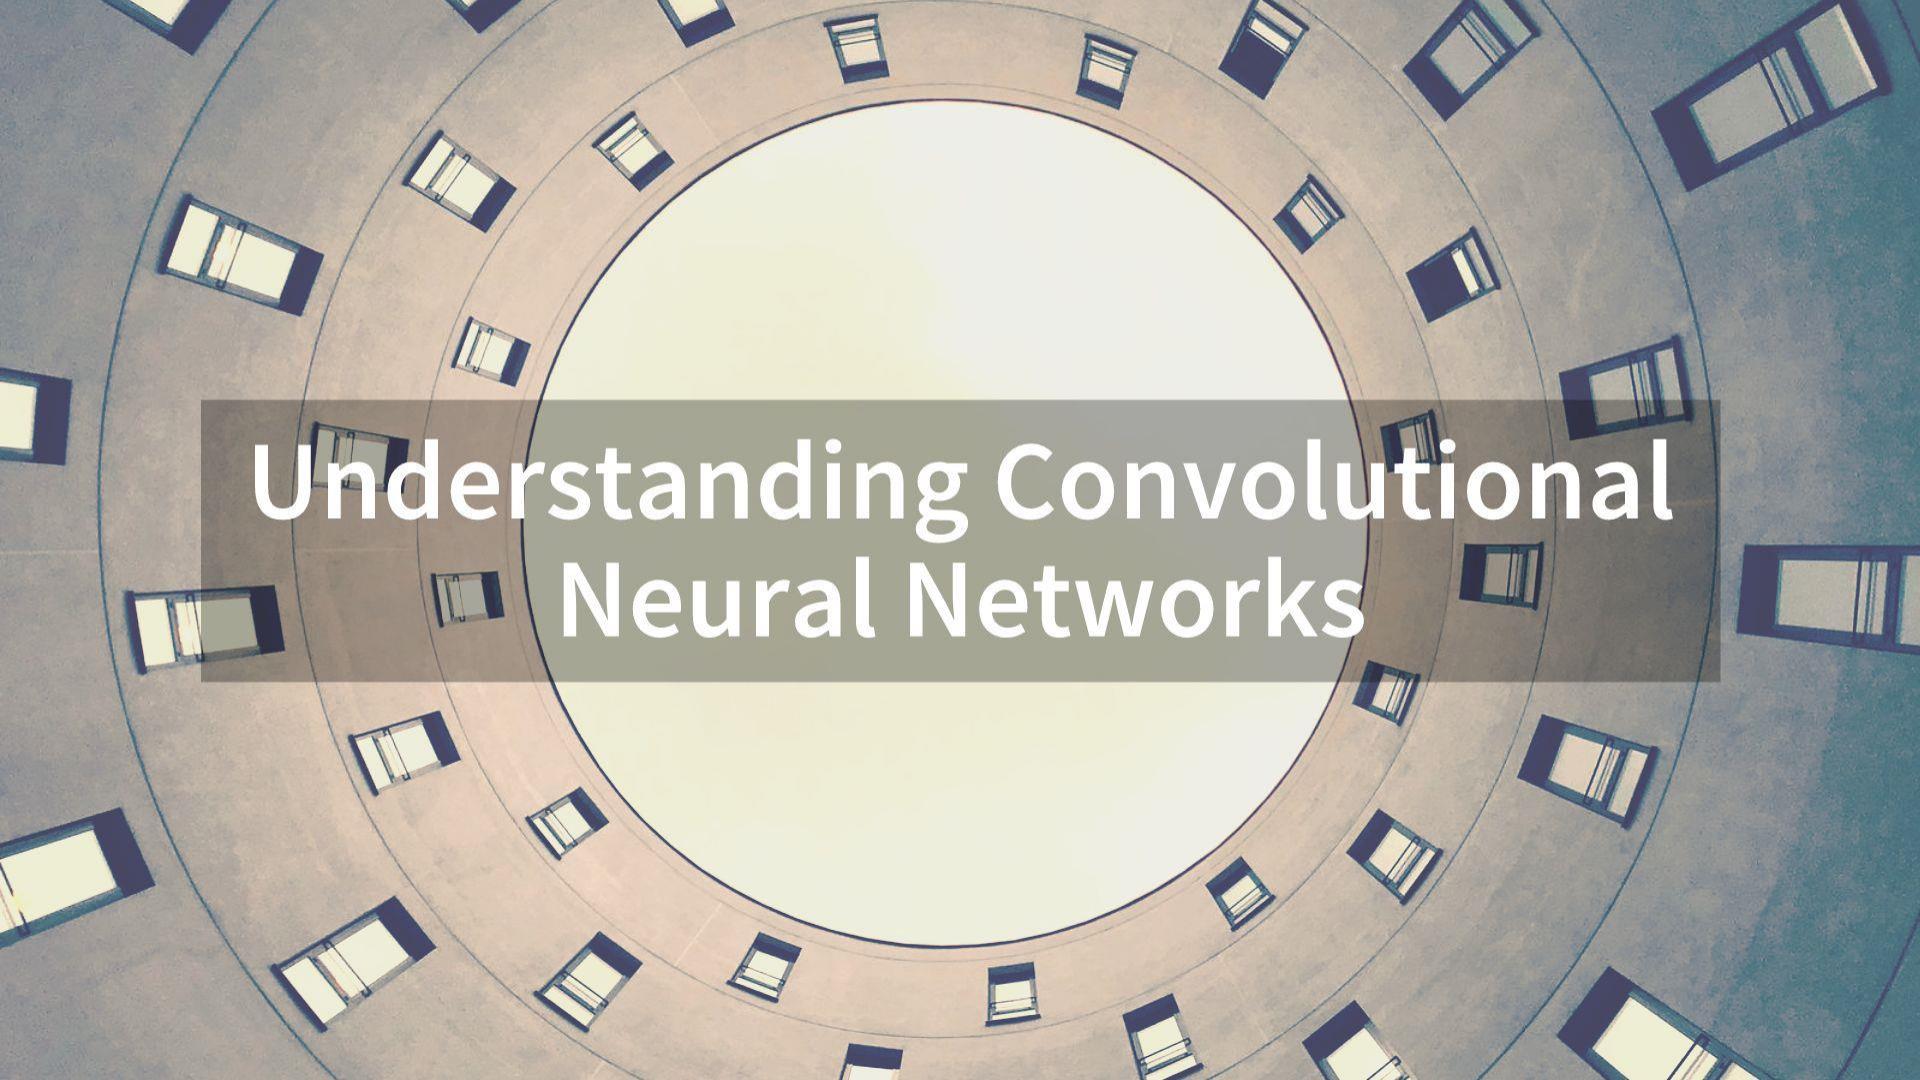


APPLIED  
MATERIALS®

 Springboard

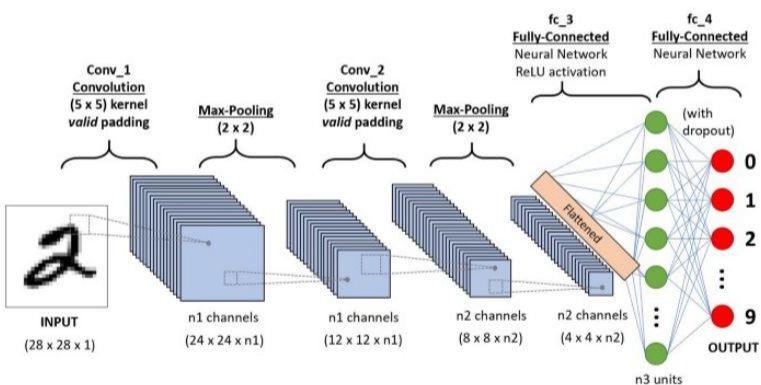


 Experts  
Machine Learning



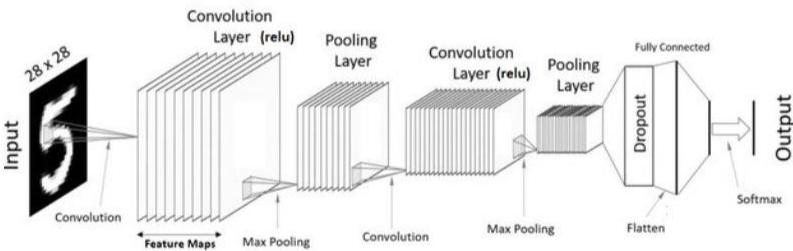
# Understanding Convolutional Neural Networks

# Convolutional Neural Networks (CNNs)



- CNNs have a layered architecture of several layers to learn hierarchical spatio-temporal features
- Convolution Layers use convolution filters to build feature maps (feature extraction)
- Pooling Layers help in reducing dimensionality after convolutions (compression)
- Non-linear activation functions are applied in the network as usual
- Dropout or BatchNormalization Layers may be used to prevent model overfitting
- FC Layers in final stages help with flattening and prediction

# CNN - Main Layer Components



- CNNs have a stacked layered architecture of several convolution and pooling layers

- **Convolution layer**

- Consists of several filters or kernels
- Passed over the entire image in patches and computes a dot product
- Result is summed up into one number per operation (dot product)

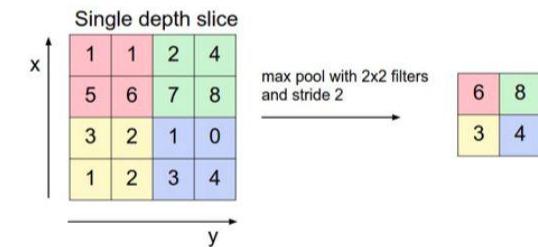
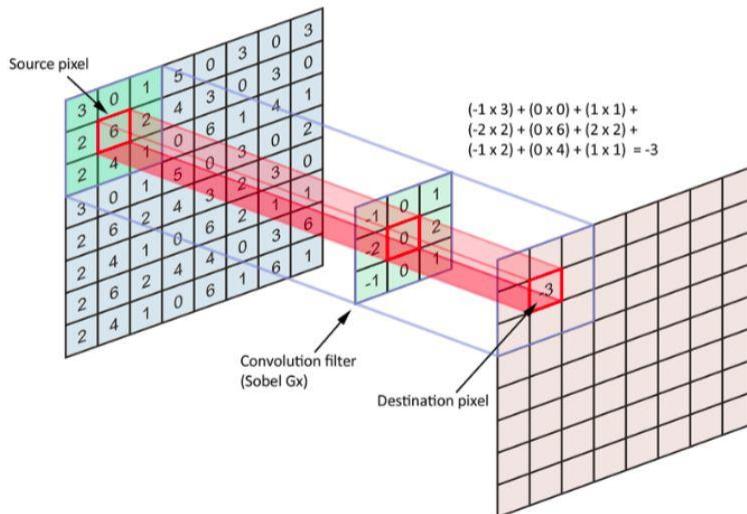
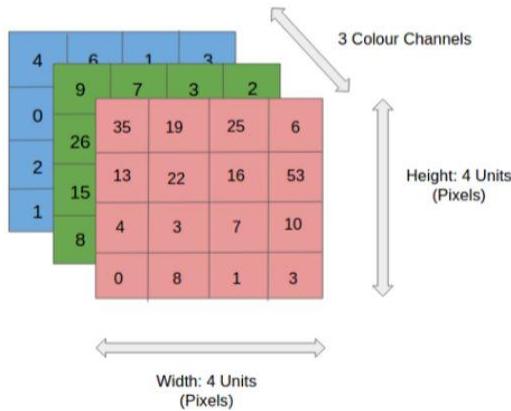
- **Pooling layer**

- Downsamples feature maps from conv layers
- Typically max-pooling is used which selects the max-pixel value out of a patch of pixels

- **Activation layer**

- Feature maps \ pooled outputs are sent through non-linear activations
- Introduces non-linearity and helps train via. backpropagation

# CNN - Conv - Pool Operations

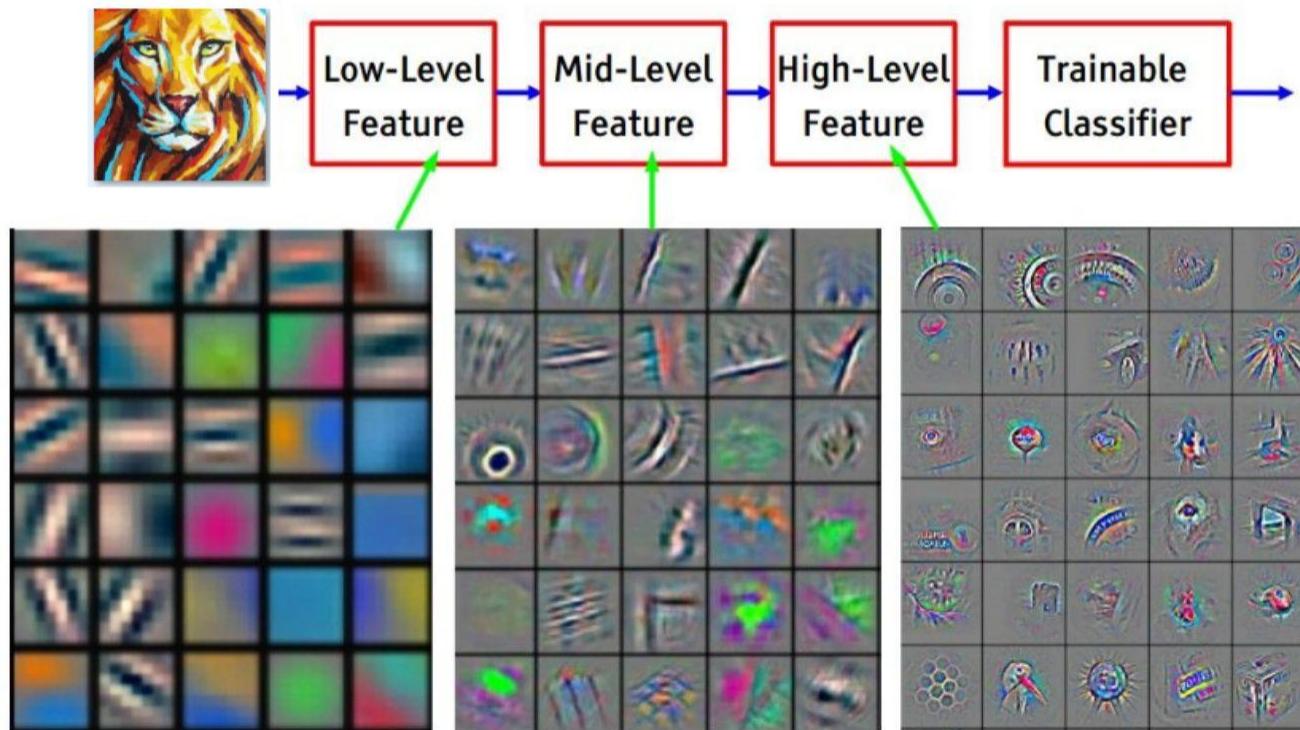


**Source Image**

**Convolution Layer**

**Pooling Layer**

# CNN - Why several layers & filters?



# CNN Layer Operations

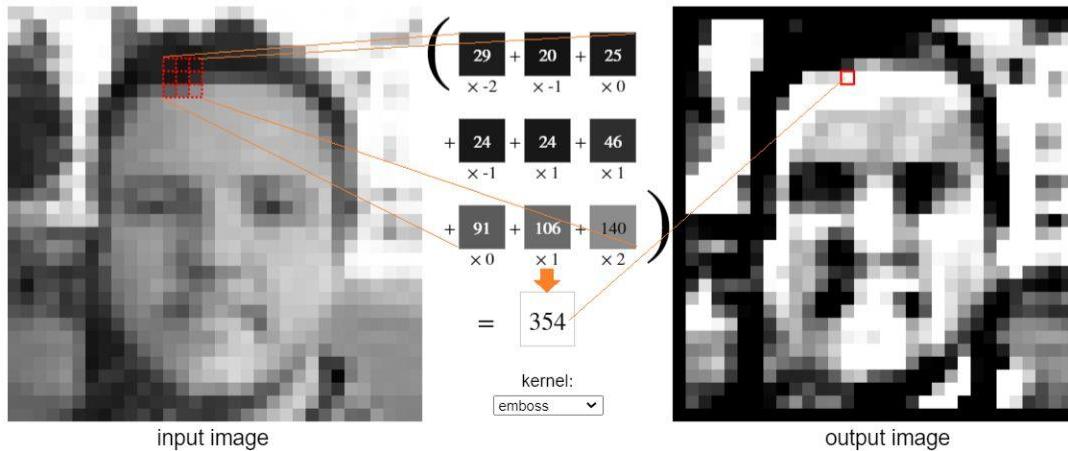
# Convolution Layer - What is a convolution operation?

emboss

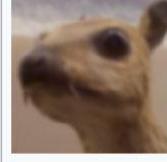
Source: <https://setosa.io/ev/image-kernels/>

$$\begin{pmatrix} -2 & -1 & 0 \\ -1 & 1 & 1 \\ 0 & 1 & 2 \end{pmatrix}$$

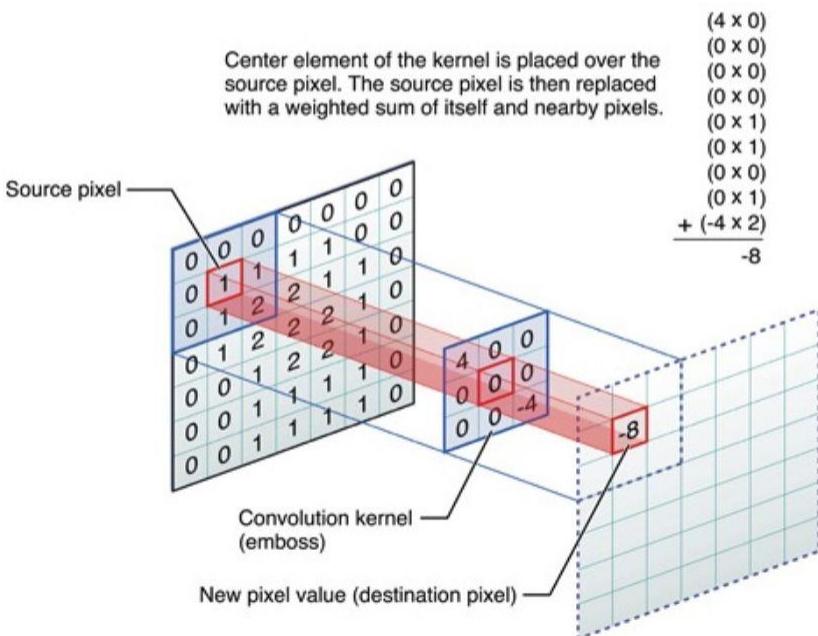
Below, for each 3x3 block of pixels in the image on the left, we multiply each pixel by the corresponding entry of the kernel and then take the sum. That sum becomes a new pixel in the image on the right.



# Convolution Layer - Types of Kernels

Operation	Kernel $\omega$	Image result $g(x,y)$	Operation	Kernel $\omega$	Image result $g(x,y)$
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$		Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$		Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix}$		Gaussian blur $3 \times 3$ (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$		Gaussian blur $5 \times 5$ (approximation)	$\frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$	

# Convolution Layer - Key Aspects



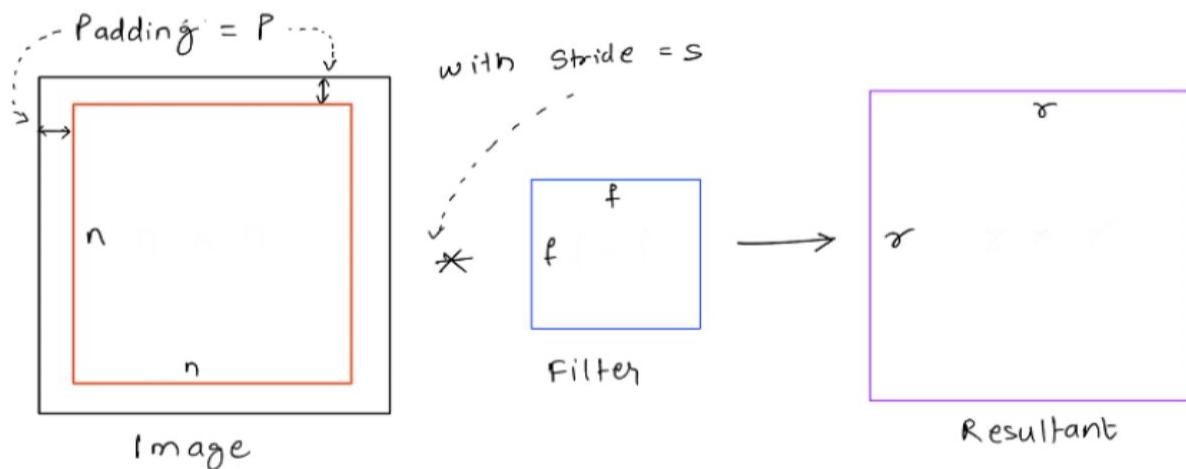
- **Kernel is also known as a filter or a feature detector**
  - Can be 3x3, 5x5 or higher filled with random or specific values based on type of filter
  - Wide variety of filters help in detecting a wide variety of features like edges, corners, textures, patterns etc.
  - Conv layers usually have multiple filters to extract multiple feature maps
- **Feature map is the output of the convolution operation**
  - Conv2d(input image, filter) gives us the convolved features
  - Elementwise multiplication of the two matrices
  - Works on the image patch-by-patch
  - Kernel size is the size of the filter applied on the image patch
  - Stride is the number of pixels to move horizontally and vertically
- **Activation layer**
  - Feature maps are sent through non-linear activations
  - Introduces non-linearity and helps train via. backpropagation

# Convolution Layer - Key Aspects

```
tf.keras.layers.Conv2D(  
    filters, kernel_size, strides=(1, 1), padding='valid', data_format=None,  
    dilation_rate=(1, 1), groups=1, activation=None, use_bias=True,  
    kernel_initializer='glorot_uniform', bias_initializer='zeros',  
    kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,  
    kernel_constraint=None, bias_constraint=None, **kwargs  
)
```

- **Filters:** Number of filters \ kernels to use
- **Kernel Size:** Size of the filter \ kernel e.g 3x3
- **Activation:** Type of non-linear activation to use e.g, relu
- **Stride:** Number of pixels to move horizontally & vertically after each convolution operation
- **Padding:** Valid or Same
  - **Valid:** reduces the size of the image after convolving due to loss of the edges
  - **Same:** pads extra pixels around the image to retain original image dimensions after convolution

# Simple Convolution - Operations on 1D Tensors



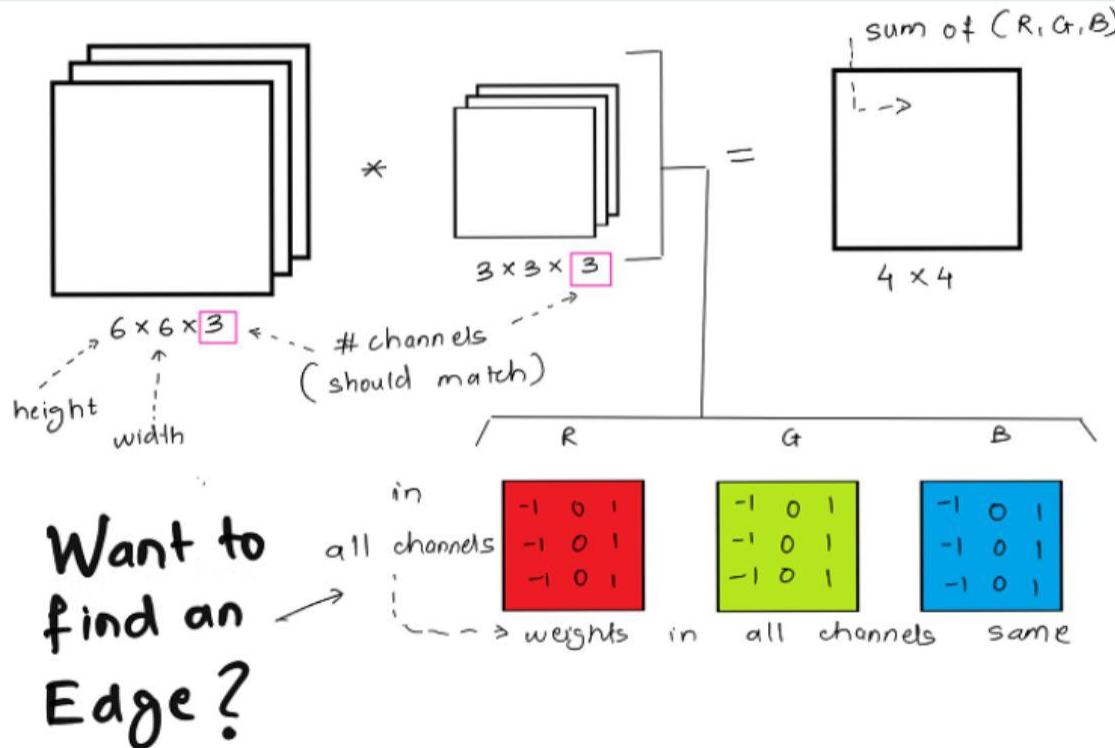
Source: <https://towardsdatascience.com/cnn-part-i-9ec412a14cb1>

A 5x5 convolutional feature map with colored cells (yellow, green, blue) and numerical values. A red box highlights the value '4' in the second row, third column. The map is labeled 'Convolved Feature'.

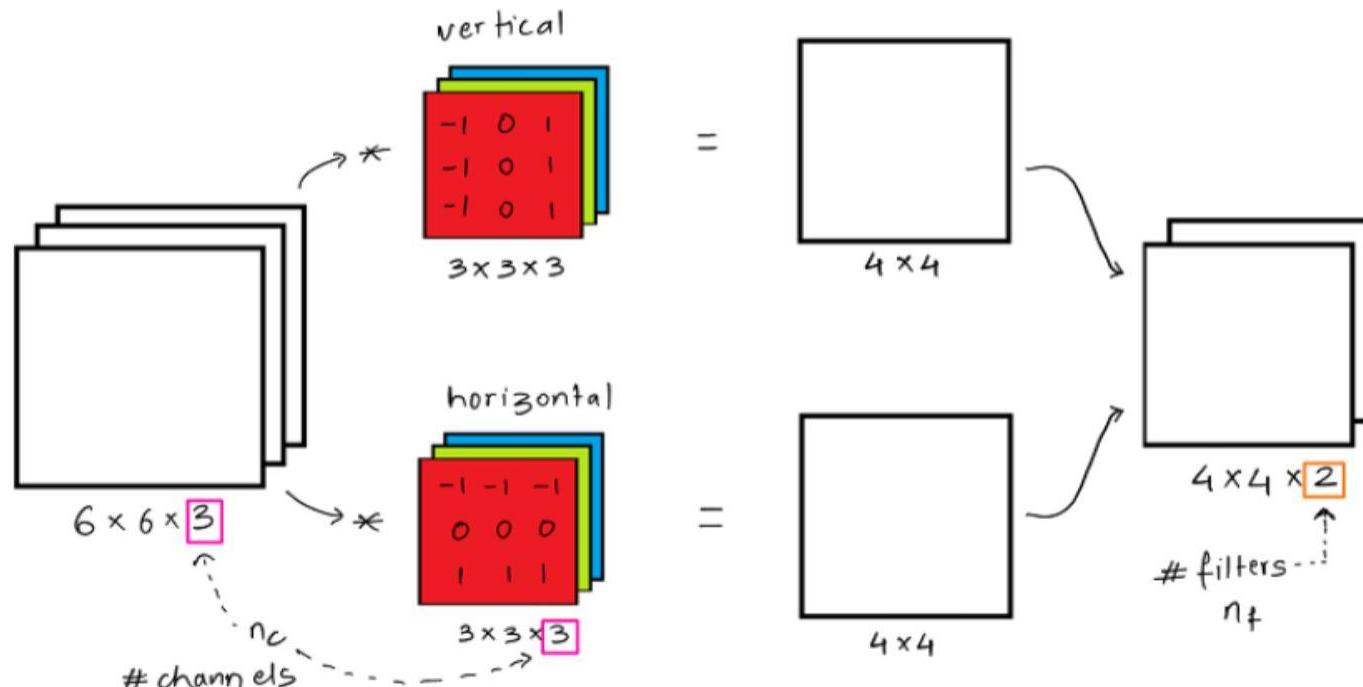
1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Convolved Feature

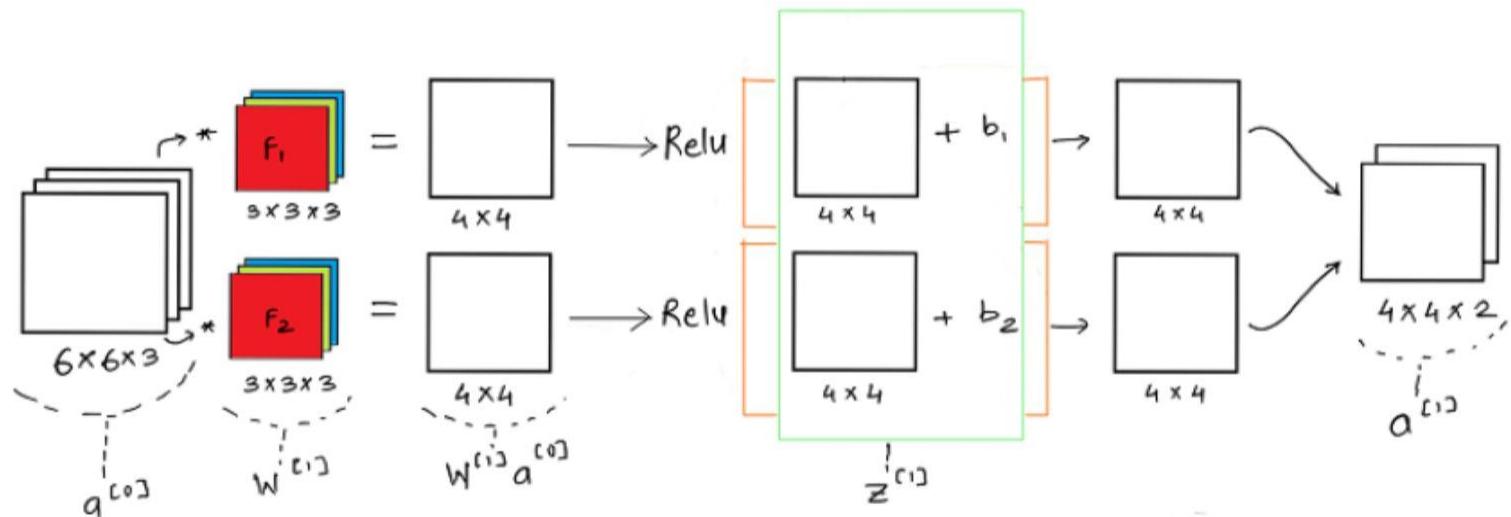
# Volume Convolution - Operations on 3D Tensors



# Volume Convolution - Multiple Filters



# Convolution Layer - Operations Flow



In any given neural network

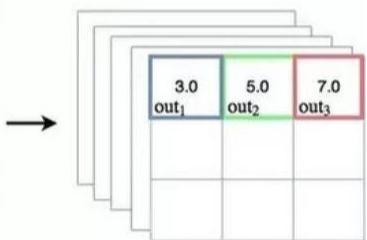
$$z^{[1]} = w^{[1]} a^{[0]} + b^{[1]}$$

$$a^{[1]} = g(z^{[1]})$$

ReLU function

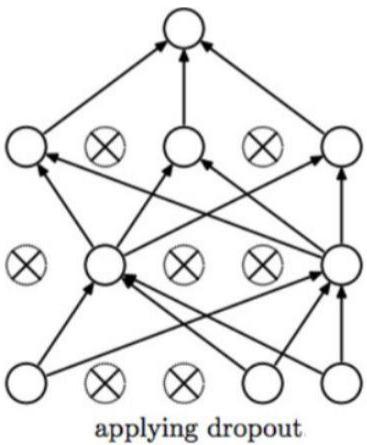
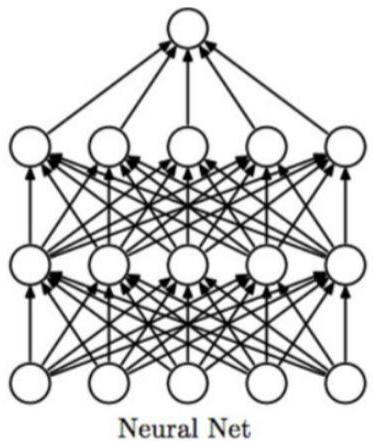
# Pooling Layer - Key Aspects

2.0 in <sub>1</sub>	3.0 in <sub>2</sub>	0.0 in <sub>3</sub>	5.0 in <sub>4</sub>	2.5 in <sub>5</sub>	0.0 pad <sub>1</sub>
2.0 in <sub>6</sub>	1.5 in <sub>7</sub>	0.5 in <sub>8</sub>	0.0 in <sub>9</sub>	7.0 in <sub>10</sub>	0.0 pad <sub>2</sub>
1.5	5.0	5.0	3.0	2.0	0.0
3.0	5.0	7.0	1.5	0.0	0.0
2.0	5.0	2.0	1.5	2.0	0.0
0.0	0.0	0.0	0.0	0.0	0.0



- Helps in reducing the number of parameters by down-sampling
- Reduction in parameters helps in faster training and prevent overfitting
- Various types of pooling are available like sum, mean, max
  - Max pooling is preferred as it is faster and usually gives a better performance than mean or sum pooling
  - Helps enhance key aspects of feature maps
- Pool size is the number of image pixels to consider for downsampling e.g, 2x2

# Dropout Layer - Key Aspects



- Technique where randomly selected neurons are ignored during training
- Helps in preventing overfitting as a regularization mechanism

A photograph showing a person's hands typing on a white Apple keyboard. In the background, a silver iMac monitor is visible, featuring its characteristic rounded design and the Apple logo. The overall scene suggests a typical office or home workspace.

# CNN Layer Operations Hands-on Examples

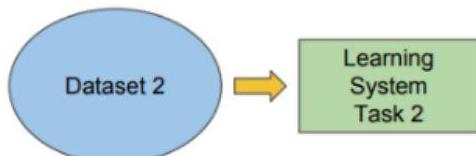
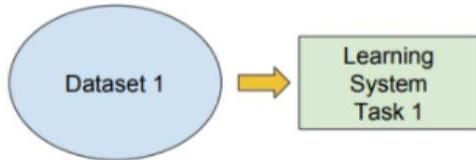
# Transfer Learning Brief



# Traditional ML vs. Transfer Learning

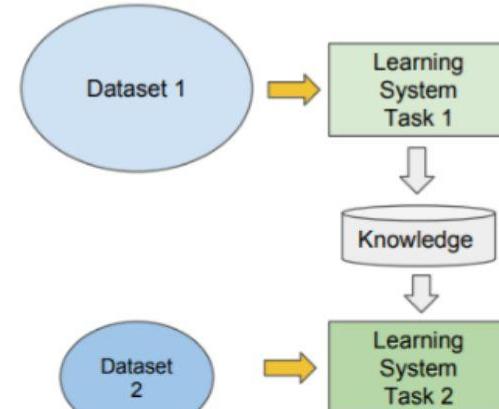
## Traditional ML

- Isolated, single task learning:
  - Knowledge is not retained or accumulated. Learning is performed w.o. considering past learned knowledge in other tasks



## vs Transfer Learning

- Learning of a new tasks relies on the previous learned tasks:
  - Learning process can be faster, more accurate and/or need less training data



# The power of Transfer Learning

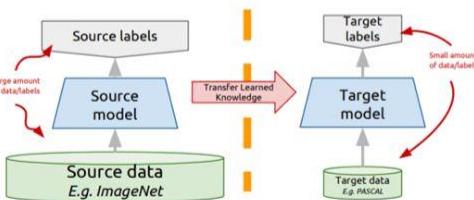
## Transfer learning: idea

Instead of training a deep network from scratch for your task:

- Take a network trained on a different domain for a different **source task**
- Adapt it for your domain and your **target task**

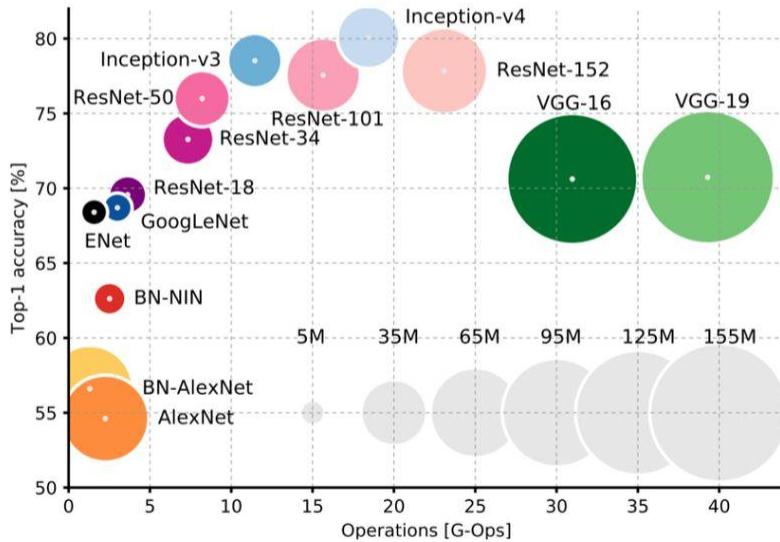
Variations:

- Same domain, different task
- Different domain, same task



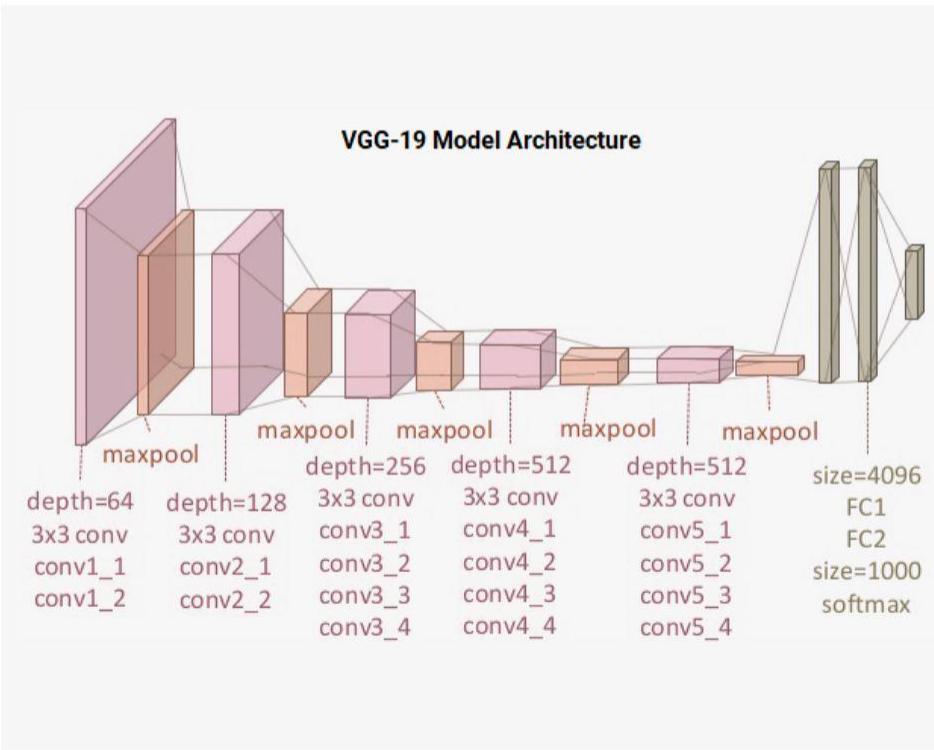
- Leverage a pre-trained deep learning model (which was trained on a large dataset — like ImageNet)
- Solve the problem of apparel classification in Fashion-MNIST by applying and transferring its knowledge in the context of our problem
- Two approaches
  - Frozen pre-trained model as a feature extractor
  - Fine-tuning pre-trained model on our data

# Transfer Learning - Advantage of Pre-trained Models



- Training DL models from scratch can take a lot of time
- Pre-trained models have been trained on a lot of data and specialized for performing specific tasks
- Can be used for both feature extraction and inference
- Models are available for computer vision and natural language processing tasks

# Pre-trained CNN - VGG-19



- VGG-19 model is a 19-layer (convolution and fully connected) deep learning network built on the ImageNet database
- 16 convolution layers using 3 x 3 convolution filters along with max pooling layers
- 2 fully connected hidden layers of 4096 units in each layer followed by a dense layer of 1000 units
- We focus on intermediate layers for extracting the right representations

A photograph of a person's hands typing on a white Apple keyboard. The keyboard is positioned in front of a light-colored Mac monitor, which has an Apple logo on its bezel. The background is a warm, indoor setting.

# CNN Classifiers Hands-on Examples



# CNN Applications

# CNN Applications

## 1 Classification & Regression

Used extensively in computer vision problems with image, video. Can also be used for audio and text

## 2 Object Detection

Can be used to detect objects in images and video

## 3 Image Segmentation

Useful for segmenting key entities in an image

## 4 Speech Processing & Modeling

Can be used for processing and modeling on speech e.g speech recognition

## 5 Style Transfer

Generative learning to transfer the style of an image on the content of an image

## 6 Search & Retrieval

Learns efficient feature representations which can be used for similarity analysis & search

## 7 Image Super Resolution

Improves the resolution of a base image with poor resolution

## 8 Image Captioning

CNN coupled with a generative language model can be used for image caption generation

# References

## ● Visuals & Content

- <http://cs231n.stanford.edu/>
- <https://www.youtube.com/playlist?list=PLkDaE6sCZn6Gl29AoE31iwdVwSG-KnDzF>
- <https://www.kaggle.com/kanncaa1/convolutional-neural-network-cnn-tutorial>
- <https://setosa.io/ev/image-kernels/>
- [https://en.wikipedia.org/wiki/Kernel\\_\(image\\_processing\)](https://en.wikipedia.org/wiki/Kernel_(image_processing))
- <https://aishack.in/tutorials/image-convolution-examples/>
- <https://towardsdatascience.com/cnn-part-i-9ec412a14cb1>
- <https://towardsdatascience.com/convolutional-neural-network-ii-a11303f807dc>
- <https://www.coursera.org/learn/introduction-tensorflow>

# Stay in Touch!



LinkedIn

<https://www.linkedin.com/in/dipanzan>



GitHub

<https://github.com/dipanjanS>



Medium

<https://medium.com/@dipanzan.sarkar>